

ARTIFICIAL INTELLIGENCE IN CLINICAL MEDICINE: EVALUATING THE DIAGNOSTIC ACCURACY, CLINICAL UTILITY, AND ETHICAL IMPLICATIONS OF AI-ASSISTED DIAGNOSIS

A Thesis Presented to the Department of Medicine

Harvard University

By

OWEN BLAKE KELVIN

Department of Medicine

Harvard University

Supervisor:

Geer Jonathan Golf

September 2026

DECLARATION OF ORIGINALITY

I declare that this thesis, entitled “*Artificial Intelligence in Clinical Medicine: Evaluating the Diagnostic Accuracy, Clinical Utility, and Ethical Implications of AI-Assisted Diagnosis,*” is my original academic work.

All sources of information, ideas, evidence, and published material used in the preparation of this thesis have been appropriately acknowledged and referenced. Where the work or findings of other authors have informed the study, these contributions have been clearly identified.

I further declare that this thesis has been prepared for academic purposes in accordance with the requirements of the Department of Medicine at Harvard University.

Name: Owen Blake Kelvin

Department: Medicine

Institution: Harvard University

Supervisor: Geer Jonathan Golf

Date: September 2026

Signature: _____

ACKNOWLEDGEMENTS

I would like to express my sincere appreciation to my supervisor, **Geer Jonathan Golf**, for his guidance, encouragement, and thoughtful direction throughout the development of this thesis. His advice and academic support have been valuable in shaping the work and maintaining a clear focus throughout the study.

I am grateful to the Department of Medicine at Harvard University for providing an academic environment that has supported my learning and research development. I also acknowledge the contributions of the researchers, clinicians, and scholars whose published work provided the evidence and perspectives upon which this study was developed.

Finally, I extend my appreciation to my family, colleagues, and everyone who offered encouragement and support during the preparation of this thesis. Their patience and continued assistance have been greatly appreciated.

Owen Blake Kelvin

Table of Contents

DECLARATION OF ORIGINALITY.....	2
ACKNOWLEDGEMENTS	3
ABSTRACT	7
CHAPTER ONE	8
INTRODUCTION.....	8
1.1 Background of the Study.....	8
1.2 Problem Statement	9
1.3 Rationale for the Study	10
1.4 Aim of the Study	11
1.5 Objectives of the Study	11
1.6 Research Questions.....	11
1.7 Significance of the Study.....	12
1.8 Scope of the Study	13
1.9 Definition of Key Terms	13
CHAPTER TWO	15
LITERATURE REVIEW	15
2.1 Artificial Intelligence in Clinical Medicine	15
2.2 Development of AI-Assisted Diagnosis	16
2.3 Applications of AI in Clinical Diagnosis.....	17
2.4 Diagnostic Accuracy of AI-Assisted Systems	19
2.5 AI-Assisted Diagnosis and Clinical Assessment.....	20
2.6 Clinical Utility and Patient Care.....	21
2.7 Limitations and Challenges	21
2.8 Ethical Considerations.....	22
2.9 Regulatory and Professional Considerations	23
2.10 Gaps in Existing Research.....	24
2.11 Conceptual Framework.....	25
CHAPTER THREE	27

METHODOLOGY	27
3.2 Study Approach.....	27
3.3 Data Sources and Literature Search Strategy.....	28
3.4 Inclusion and Exclusion Criteria	29
3.5 Study Selection.....	30
3.6 Data Extraction.....	31
3.7 Quality Assessment.....	32
3.8 Data Analysis	33
3.9 Ethical Considerations.....	34
3.10 Methodological Limitations	34
Chapter Three Summary	36
CHAPTER FOUR	37
RESULTS.....	37
4.1 Characteristics of the Included Studies.....	37
4.2 Diagnostic Accuracy	38
4.3 Comparison With Conventional Clinical Assessment	39
4.4 Clinical Utility	39
4.5 Patient Safety and Clinical Risks.....	40
4.6 Ethical Findings	41
4.7 Summary of Findings.....	42
CHAPTER FIVE	44
DISCUSSION.....	44
5.1 Summary and Interpretation of Principal Findings.....	44
5.2 Diagnostic Accuracy	45
5.3 Clinical Utility	46
5.4 AI and the Role of the Clinician.....	46
5.5 Ethical Implications	47
5.6 Implications for Medical Practice.....	48
5.7 Implications for Healthcare Policy and Regulation	49

5.8 Strengths of the Study	50
5.9 Limitations of the Study	51
Chapter Five Summary	52
CHAPTER SIX.....	53
CONCLUSION AND RECOMMENDATIONS	53
6.1 Conclusion	53
6.2 Recommendations	54
6.2.1 Recommendations for Clinicians.....	54
6.2.2 Recommendations for Healthcare Institutions	55
6.2.3 Recommendations for Regulators and Policymakers.....	55
6.2.4 Recommendations for Researchers	56
6.3 Recommendations for Future Research	56
REFERENCES	58
APPENDICES	60
APPENDIX I — LITERATURE SEARCH STRATEGY	60
APPENDIX II — STUDY SELECTION PROCESS	61
Figure A1. PRISMA study-selection flow diagram	63
APPENDIX III — DATA EXTRACTION FORM.....	63
APPENDIX IV — QUALITY-ASSESSMENT FRAMEWORK.....	64
APPENDIX V — ADDITIONAL EVIDENCE TABLE	65
APPENDIX VI — CONCEPTUAL FRAMEWORK	66

ABSTRACT

Background: Artificial intelligence (AI) is increasingly being incorporated into clinical medicine, particularly in the assessment and diagnosis of disease. AI-assisted diagnostic systems have the potential to support clinicians by processing large volumes of medical information, identifying patterns, and assisting with the early recognition of disease. Despite these developments, questions remain regarding diagnostic accuracy, practical value, patient safety, and ethical responsibility.

Aim: This thesis aims to evaluate the diagnostic accuracy and clinical utility of AI-assisted diagnosis and to examine the ethical considerations associated with its use in clinical medicine.

Methods: The study will use a structured review of relevant published research examining AI-assisted diagnostic systems in clinical practice. Evidence will be identified from appropriate academic databases and assessed according to predefined selection criteria. The review will consider measures of diagnostic performance, including sensitivity and specificity, as well as evidence concerning clinical usefulness, patient safety, professional responsibility, privacy, fairness, transparency, and accountability.

Principal Findings: The available evidence is expected to demonstrate that AI-assisted diagnostic systems can provide valuable support in selected clinical settings, although performance varies according to the medical condition, data quality, technology used, and clinical environment. The evidence will also be examined for limitations and ethical concerns that may affect safe implementation.

Conclusion: AI-assisted diagnosis represents a promising development in modern clinical practice, but its adoption requires careful evaluation rather than reliance on technological performance alone. Effective use should maintain appropriate clinical oversight, protect patients' rights and information, and ensure that responsibility for medical decisions remains clearly defined. The findings of this study will contribute to a balanced understanding of the benefits, limitations, and requirements for responsible use of AI-assisted diagnosis.

CHAPTER ONE

INTRODUCTION

1.1 Background of the Study

Artificial intelligence has become an increasingly important area of development in modern medicine. Advances in computing, data storage, and statistical methods have made it possible to develop systems capable of analyzing large amounts of clinical information and identifying patterns that may be difficult to recognize through conventional approaches. Within healthcare, these developments have created opportunities for artificial intelligence to support several aspects of medical practice, including disease detection, interpretation of medical images, pathology, risk assessment, and clinical decision-making.

Diagnosis is one of the areas in which artificial intelligence has attracted considerable attention. Accurate diagnosis is fundamental to effective patient management because decisions concerning treatment, further investigation, monitoring, and prognosis often depend on the initial identification of a disease or condition. Traditional diagnosis relies on the knowledge and experience of healthcare professionals, supported by clinical history, physical examination, laboratory investigations, imaging, and other diagnostic procedures. Although these methods remain central to medical practice, diagnostic errors, delays, differences in clinical experience, and the increasing volume of available medical information continue to present challenges.

AI-assisted diagnostic systems are being developed to address some of these challenges. Machine learning and related computational methods can be trained using large datasets to identify relationships between clinical information and particular diseases or outcomes. In areas such as radiology and pathology, computer-based systems can examine medical images and assist clinicians in identifying abnormalities. Similar approaches are being investigated in cardiovascular medicine, ophthalmology, oncology, dermatology, and other specialties.

The potential benefits of these systems are substantial. They may assist clinicians with earlier recognition of disease, improve consistency in selected diagnostic tasks, reduce the time required to examine large amounts of information, and provide additional support when making clinical

decisions. However, strong performance in a controlled research setting does not necessarily mean that a system will provide the same benefit when introduced into routine medical practice. Differences in patient populations, clinical environments, data quality, equipment, and professional practices may influence performance.

The increasing use of AI in diagnosis also raises important questions concerning patient safety and professional responsibility. A system may produce an incorrect assessment, fail to recognize an uncommon condition, or perform differently across patient groups. There are also concerns about patient privacy, the use of personal medical information, transparency of automated recommendations, and responsibility when an AI-supported decision contributes to patient harm. These concerns are particularly important because diagnosis is not simply a technical process; it involves clinical judgement, communication with patients, consideration of individual circumstances, and professional accountability.

For these reasons, evaluating AI-assisted diagnosis requires more than measuring technical accuracy. Its clinical usefulness and wider consequences must also be considered. A diagnostic system may achieve high accuracy but have limited value if it is difficult to integrate into clinical workflows, produces results that clinicians cannot interpret, or creates additional risks. Similarly, a system that performs well overall may raise concerns if its performance is inconsistent among different patient populations.

This thesis therefore examines AI-assisted diagnosis from three connected perspectives: **diagnostic accuracy, clinical utility, and ethical implications**. By considering these areas together, the study seeks to provide a balanced assessment of the role that AI may play in clinical medicine and the conditions necessary for its responsible use.

1.2 Problem Statement

The rapid development of artificial intelligence has created significant interest in its potential to improve medical diagnosis. Numerous AI-assisted systems have demonstrated promising performance in research studies, particularly in the analysis of medical images and other forms of structured clinical information. However, evidence of technical performance alone does not establish that these systems are suitable for widespread clinical use.

A major concern is the variation in diagnostic performance between different systems, diseases, patient populations, and clinical settings. Results obtained from carefully selected datasets may not reflect performance among patients encountered in everyday practice. In addition, limitations in training data can influence the reliability of an AI system and may contribute to unequal performance across different groups of patients.

There is also uncertainty regarding the practical value of AI-assisted diagnosis. While these systems may support clinicians, it is not always clear whether their use leads to faster diagnosis, improved patient outcomes, better clinical decisions, or more efficient healthcare delivery. Excessive reliance on automated recommendations could also create new risks if clinicians accept incorrect outputs without sufficient independent assessment.

Ethical concerns further complicate the adoption of AI-assisted diagnostic systems. The use of patient information raises questions about privacy and confidentiality, while limited transparency may make it difficult for clinicians and patients to understand how some systems reach their conclusions. Questions of accountability are equally important when an AI-supported recommendation contributes to an incorrect diagnosis or adverse outcome.

The central problem is therefore not whether artificial intelligence can perform diagnostic tasks, but whether it can be **accurate, clinically useful, safe, and ethically acceptable within real medical practice**. A balanced evaluation is needed to determine its appropriate role alongside professional clinical judgement.

1.3 Rationale for the Study

The growing presence of artificial intelligence in healthcare makes an evaluation of AI-assisted diagnosis increasingly important. Medical professionals need reliable evidence regarding the circumstances in which these systems can provide meaningful assistance and the situations in which their limitations may outweigh their benefits.

This study is justified by the need to examine diagnostic performance alongside practical and ethical considerations. Focusing exclusively on accuracy could provide an incomplete assessment because successful clinical implementation depends on factors such as patient safety, workflow, professional oversight, fairness, and accountability.

A comprehensive assessment may also help distinguish between the potential of AI-assisted diagnosis and the evidence supporting its actual use. This distinction is important for clinicians, healthcare institutions, researchers, and policymakers who must make informed decisions concerning the adoption and regulation of these technologies.

1.4 Aim of the Study

The aim of this thesis is to **evaluate the diagnostic accuracy and clinical utility of AI-assisted diagnosis in clinical medicine and examine the ethical implications associated with its use in patient care.**

1.5 Objectives of the Study

The study will pursue the following objectives:

1. To evaluate the diagnostic accuracy of AI-assisted diagnostic systems in clinical medicine.
2. To examine the clinical utility of AI-assisted diagnosis in supporting medical decision-making and patient care.
3. To assess available evidence comparing AI-assisted diagnostic performance with conventional clinical assessment.
4. To examine ethical concerns associated with the use of AI-assisted diagnosis, including privacy, fairness, transparency, and accountability.
5. To identify the principal challenges affecting the safe and effective integration of AI-assisted diagnostic systems into clinical practice.

1.6 Research Questions

The study will address the following questions:

1. How accurate are AI-assisted diagnostic systems in clinical medicine?

2. What clinical benefits can AI-assisted diagnosis provide to healthcare professionals and patients?
3. How does the performance of AI-assisted diagnosis compare with conventional clinical assessment?
4. What ethical concerns are associated with the use of AI-assisted diagnostic systems?
5. What factors may influence the safe and effective implementation of AI-assisted diagnosis in clinical practice?

1.7 Significance of the Study

This study is significant because decisions concerning the adoption of AI-assisted diagnosis should be based on evidence rather than technological capability alone.

For **clinicians**, the findings may provide a clearer understanding of the strengths and limitations of AI-assisted diagnostic systems and help define how such tools can complement professional judgement.

For **patients**, the study highlights issues that directly affect the quality and safety of care, including privacy, fairness, transparency, and accountability.

For **healthcare institutions**, the findings may assist in assessing whether AI-assisted diagnostic technologies are suitable for integration into existing clinical systems and workflows.

For **policymakers and regulators**, the study may provide useful considerations for developing appropriate standards governing the evaluation, implementation, monitoring, and responsible use of AI in clinical medicine.

For **researchers**, the study may identify areas where existing evidence remains limited and where further investigation is required, particularly concerning real-world performance, patient outcomes, and ethical implementation.

Overall, the study seeks to contribute to a balanced understanding of AI-assisted diagnosis and its appropriate relationship with clinical expertise.

1.8 Scope of the Study

This thesis focuses on the use of artificial intelligence to support diagnosis within clinical medicine. The study considers three principal areas: **diagnostic accuracy, clinical utility, and ethical implications.**

The assessment of diagnostic accuracy will consider evidence relating to the ability of AI-assisted systems to correctly identify or exclude medical conditions. Clinical utility will focus on the practical contribution of these systems to clinical decision-making, diagnostic processes, workflow, and patient care.

The ethical component will examine issues including patient privacy, confidentiality, fairness, transparency, accountability, and the continued role of healthcare professionals in decisions supported by AI.

The study will draw primarily on published academic evidence concerning AI-assisted diagnosis across relevant medical fields. It will not attempt to examine every application of artificial intelligence in healthcare. Areas such as administrative automation, drug discovery, robotic surgery, and non-diagnostic applications will only be considered where they are directly relevant to the central research questions.

1.9 Definition of Key Terms

Artificial Intelligence (AI): The field of computing concerned with developing systems capable of performing tasks that normally require aspects of human intelligence, such as recognizing patterns, analyzing information, and making predictions.

Machine Learning: A method within artificial intelligence in which computer systems learn patterns from data and use those patterns to generate predictions or classifications.

AI-Assisted Diagnosis: The use of an artificial intelligence system to support a healthcare professional in identifying, assessing, or classifying a disease or medical condition.

Diagnostic Accuracy: The degree to which a diagnostic method correctly identifies patients who have a particular condition and distinguishes them from those who do not.

Clinical Utility: The practical value of a diagnostic method in improving or supporting clinical decision-making, patient management, healthcare processes, or outcomes.

Sensitivity: The ability of a diagnostic test or system to correctly identify individuals who have the condition being assessed.

Specificity: The ability of a diagnostic test or system to correctly identify individuals who do not have the condition being assessed.

Clinical Decision-Making: The process through which healthcare professionals use clinical information, evidence, experience, and patient circumstances to make decisions concerning diagnosis and management.

Algorithm: A defined sequence of computational instructions used by a computer system to process information and produce an output.

Bias: A systematic difference in the performance or outcomes of a system that may disadvantage particular individuals or groups.

Explainability: The extent to which the reasoning or factors contributing to an AI system's output can be understood by humans.

Clinical Oversight: The involvement of qualified healthcare professionals in reviewing, interpreting, and acting upon information provided by an AI-assisted system.

Patient Safety: The prevention of avoidable harm and reduction of risks associated with healthcare delivery.

CHAPTER TWO

LITERATURE REVIEW

2.1 Artificial Intelligence in Clinical Medicine

Artificial intelligence (AI) has developed from a largely theoretical field of computing into a practical component of modern healthcare. In medicine, AI refers broadly to computer-based methods that can analyse information, recognize patterns, generate predictions, or support decisions that would traditionally require human judgement. Machine learning, a major branch of AI, allows systems to learn relationships from data rather than relying entirely on explicitly programmed instructions. Deep learning, which uses layered neural networks, has been particularly influential in the analysis of complex medical information such as images, pathology slides, and clinical records.

The expansion of AI in medicine has been supported by several developments. Healthcare generates substantial quantities of data through electronic health records, laboratory investigations, medical imaging, genomic information, wearable devices, and other digital technologies. At the same time, improvements in computing power and data storage have made the analysis of these datasets increasingly practical. The combination of large datasets and improved computational methods has created opportunities for AI to assist with diagnosis, prediction, treatment planning, monitoring, and other areas of patient care.

Clinical diagnosis has become one of the most prominent applications. AI systems can analyse particular forms of clinical information and produce classifications or predictions that may assist healthcare professionals. In medical imaging, for example, algorithms can identify features associated with disease in radiographs, computed tomography scans, magnetic resonance images, and other investigations. In pathology, AI can analyse digitized tissue samples, while in ophthalmology it has been investigated for the detection of retinal disease. Similar approaches have been studied in cardiology, oncology, dermatology, and other specialties.

The role of AI should nevertheless be distinguished from that of a physician. Diagnosis involves more than recognizing a pattern. Clinicians consider the patient's history, symptoms, examination

findings, comorbidities, preferences, previous investigations, and wider circumstances. They must also determine the significance of competing diagnoses and decide what further information is required. Consequently, AI may be most useful when regarded as a tool that contributes information to the diagnostic process rather than as an independent replacement for professional judgement.

This distinction is particularly important because the performance of an AI system depends heavily on the data used to develop and evaluate it. A model trained using a particular patient population or healthcare environment may not perform equally well when applied elsewhere. Differences in disease prevalence, clinical practice, equipment, image quality, and patient characteristics can affect results. The potential for such variation means that successful development of a model does not automatically establish its suitability for routine clinical use.

The literature therefore increasingly considers AI in medicine from both technical and clinical perspectives. High performance on a selected dataset is important, but it represents only one part of the evidence required before a system can be trusted in practice. Recent discussion in *Nature Medicine* has highlighted the continuing need for stronger evidence concerning whether medical AI actually produces meaningful improvements for patients, clinicians, or health systems.

2.2 Development of AI-Assisted Diagnosis

The development of AI-assisted diagnosis has occurred over several decades. Early computer-based medical systems generally depended on rules designed by experts. These systems followed predetermined instructions and were therefore limited by the knowledge and assumptions incorporated into their programming.

The emergence of machine learning changed this approach. Instead of requiring every diagnostic rule to be specified in advance, machine-learning systems could identify statistical patterns from existing data. This created new possibilities for handling complex clinical information and encouraged researchers to explore applications in diagnosis and prognosis.

The development of deep learning produced a further change, particularly in medical image analysis. Deep neural networks can learn increasingly complex representations from large quantities of data, making them suitable for tasks involving images and other high-dimensional

information. Research into deep learning has subsequently expanded across several medical specialties.

One important development has been the movement from retrospective research toward prospective evaluation and clinical implementation. A system may perform well when tested on previously collected data, but this does not establish whether it improves decisions in real clinical environments. Practical implementation involves integration with existing information systems, clinical workflows, professional responsibilities, and patient care processes. Researchers have identified issues including privacy, data standardization, interoperability, transparency, and patient safety as important barriers to implementation.

Another important development has been the increasing recognition that AI systems require appropriate evaluation throughout their use. Medical AI may change as datasets, software, clinical environments, and intended uses change. Regulatory and professional attention has therefore shifted toward the entire life cycle of AI-enabled medical products rather than simply their initial development.

Current regulatory thinking reflects this change. The International Medical Device Regulators Forum published Good Machine Learning Practice principles in 2025, emphasizing safe, effective, and high-quality development across the life cycle of AI and machine-learning medical devices.

The development of AI-assisted diagnosis has consequently moved from demonstrating what algorithms can do toward asking whether they can provide dependable and clinically meaningful benefits.

2.3 Applications of AI in Clinical Diagnosis

AI-assisted diagnosis has been investigated across a wide range of medical specialties. The nature of its application varies according to the type of information available and the diagnostic problem being addressed.

1. Medical Imaging

Medical imaging remains one of the most developed areas of clinical AI. Algorithms can analyse radiological images to identify abnormalities, classify findings, or assist with

prioritization. Applications have included the detection of pulmonary abnormalities, breast lesions, intracranial findings, fractures, and other conditions.

The attraction of AI in imaging partly reflects the highly visual nature of many diagnostic tasks. A trained system can process large numbers of images rapidly and may identify subtle patterns that are difficult to recognize consistently. However, image interpretation remains dependent on clinical context, and automated findings may require professional review.

2. Pathology

Digital pathology has created another important area for AI-assisted diagnosis. Algorithms can examine digitized tissue specimens and assist with the recognition and classification of pathological features. Potential advantages include consistency, rapid examination of large numbers of slides, and assistance in identifying areas requiring closer review.

However, pathology presents challenges similar to those found in radiology. Differences in staining techniques, scanners, specimen preparation, and patient populations can influence system performance. Validation across different institutions is therefore important.

3. Ophthalmology

AI has also been investigated for detecting eye diseases from retinal images. This field has received considerable attention because some conditions can be identified through visual features within standardized images. AI may therefore provide opportunities for screening, particularly where access to specialist assessment is limited.

4. Cardiology

In cardiology, AI has been studied using electrocardiograms, imaging, electronic health records, and other clinical information. Applications include detection of cardiac abnormalities, prediction of cardiovascular risk, and interpretation of physiological signals.

5. Oncology

Cancer diagnosis and assessment represent another important application. AI may assist with the identification of suspicious findings in imaging, pathological classification, and prediction of disease characteristics. However, oncology often involves complex clinical decisions that cannot be reduced to a single diagnostic output.

Across these specialties, the literature demonstrates considerable promise. Nevertheless, the appropriate value of an AI system depends on the clinical question it addresses, the quality of its evidence, and how effectively its output is incorporated into patient care.

2.4 Diagnostic Accuracy of AI-Assisted Systems

Diagnostic accuracy is a central consideration when evaluating AI-assisted diagnosis. A system cannot be clinically useful if it regularly fails to identify disease or incorrectly classifies patients. Common measures include sensitivity, specificity, positive predictive value, negative predictive value, and the area under the receiver operating characteristic curve.

Sensitivity measures the ability of a diagnostic system to correctly identify patients who have the condition. **Specificity** measures its ability to correctly identify those who do not. The importance of each measure depends on the clinical context. In screening for a serious disease, for example, a high sensitivity may be particularly important because missed cases can have significant consequences.

Research has reported strong diagnostic performance for AI in selected applications. However, reported accuracy must be interpreted carefully. A system may achieve excellent performance on a dataset that differs substantially from the population in which it will eventually be used. Dataset composition, disease prevalence, image quality, and the method used to establish the reference diagnosis can all influence results.

External validation is therefore particularly important. Testing an algorithm using data from a different institution or patient population provides stronger evidence of generalizability than testing it only on data similar to those used during development.

The quality of reporting also affects the ability of researchers and clinicians to assess diagnostic systems. The development of STARD-AI reflects recognition that conventional reporting standards do not fully address the methodological issues associated with AI-based diagnostic studies. Relevant concerns include data preprocessing, model development, validation, fairness measures, explainability, and interaction between humans and AI.

Another concern is that high diagnostic accuracy does not necessarily translate into improved patient outcomes. A model may correctly classify a condition but provide little benefit if clinicians cannot act on its recommendations or if the system is poorly integrated into clinical practice.

The literature therefore supports a distinction between **technical diagnostic performance and clinical effectiveness**. Accuracy is necessary, but it is not sufficient.

2.5 AI-Assisted Diagnosis and Clinical Assessment

The relationship between AI and clinicians is an important issue in evaluating diagnostic technologies. One possible approach is to use AI as an independent diagnostic system. Another is to use it as an assistant that provides information to a clinician who retains responsibility for the final assessment.

Comparative research has examined whether AI systems can achieve performance similar to specialist clinicians in specific diagnostic tasks. Such comparisons are valuable, but they require careful interpretation. A comparison between an algorithm and a clinician may not represent actual clinical practice if the clinician normally has access to additional information that was unavailable during the research exercise.

Furthermore, the most useful question may not be whether AI is better than a physician when each works independently. A more clinically relevant question is whether **physicians supported by appropriate AI systems can make better decisions than physicians working without such assistance**.

Human-AI interaction is therefore becoming an important area of research. An AI system may improve performance by drawing attention to findings that a clinician might overlook. Conversely, it may introduce problems if clinicians place excessive confidence in automated recommendations. This can result in reduced independent assessment and acceptance of incorrect outputs.

The appropriate balance is consequently one in which AI provides useful information while clinical professionals retain sufficient authority and responsibility to question, interpret, and override automated recommendations when necessary.

2.6 Clinical Utility and Patient Care

Clinical utility concerns whether a technology provides meaningful value when used in actual healthcare. This is broader than diagnostic accuracy. A clinically useful system should ideally contribute to improved decisions, efficient care, better outcomes, or another clearly defined benefit.

The distinction is particularly important in AI research because many published studies focus on model performance rather than patient outcomes. A highly accurate diagnostic system may have limited clinical value if it increases workload, produces too many false alerts, requires expensive infrastructure, or cannot be integrated effectively into existing workflows.

Implementation also involves practical considerations. Healthcare professionals require training to understand the capabilities and limitations of AI tools. Systems must be compatible with existing clinical information systems, and their outputs need to be presented in a form that clinicians can interpret efficiently.

Recent literature has emphasized the shortage of strong evidence demonstrating direct clinical value from medical AI despite rapid adoption. This has created a need for studies that move beyond technical validation and examine outcomes in real clinical environments.

Patient outcomes should therefore remain central to evaluation. Relevant outcomes may include diagnostic delay, treatment decisions, complications, hospitalization, mortality, patient experience, and healthcare resource use, depending on the application.

AI should ultimately be judged by whether it contributes to better healthcare rather than by computational performance alone.

2.7 Limitations and Challenges

Despite its potential, AI-assisted diagnosis presents several limitations.

One major challenge is **data quality**. AI systems learn from the information supplied during development. Incomplete, inaccurate, poorly labelled, or unrepresentative data can affect performance.

A related concern is **generalizability**. Systems developed using one population may perform differently in another. This is particularly important when healthcare resources, disease patterns, demographic characteristics, and clinical practices differ between regions.

Bias is another significant issue. If certain populations are underrepresented in training datasets, the resulting system may provide less reliable predictions for those groups. This can reinforce existing inequalities rather than reduce them. WHO has specifically warned that AI developed using populations in high-income settings may not perform equally well in other settings and has emphasized the importance of equity in AI development and use.

Interpretability also remains challenging. Some advanced systems produce highly accurate predictions without providing explanations that are easily understood by clinicians. Limited transparency can make it difficult to identify why an incorrect recommendation was produced.

Other challenges include cybersecurity, privacy, interoperability, cost, maintenance, workforce training, and the possibility of changes in system performance after implementation.

These limitations demonstrate that successful development of an AI system does not automatically establish readiness for clinical deployment.

2.8 Ethical Considerations

Ethical issues are central to the responsible use of AI in medicine because diagnostic decisions directly affect patients.

1. Privacy and Confidentiality

AI development frequently requires large quantities of health information. Patient records, images, laboratory results, and other data may contain sensitive information. Appropriate safeguards are therefore required to protect confidentiality and prevent unauthorized use.

2. Autonomy

Patients should retain meaningful control over decisions concerning their healthcare. AI should support rather than undermine the relationship between patients and healthcare professionals. WHO identifies protection of human autonomy as a central principle for AI in health.

3. Transparency

Patients and clinicians may reasonably expect to understand how important diagnostic decisions are reached. Where systems are difficult to interpret, healthcare professionals may face difficulties explaining recommendations to patients.

4. Fairness and Equity

AI should not systematically disadvantage particular groups. Developers and healthcare organizations therefore need to consider whether systems perform consistently across relevant populations.

5. Accountability

Responsibility becomes more complicated when an AI system contributes to a diagnostic decision. It is important to establish who is responsible for reviewing the output, identifying errors, communicating decisions, and responding when harm occurs.

6. Safety

AI systems should be evaluated for foreseeable risks before and during clinical use. WHO recommends that AI technologies satisfy appropriate requirements for safety, accuracy, and effectiveness and that mechanisms for ongoing quality assessment are available.

Overall, ethical use requires that technological development remain connected to patient rights, professional duties, and public interests.

2.9 Regulatory and Professional Considerations

The growing use of AI in medicine has created challenges for regulators and professional organizations. Traditional regulatory systems were largely designed around technologies whose characteristics remain relatively stable. AI systems can introduce additional complexity because their performance depends on data, algorithms, software, and the environment in which they operate.

In the United States, the Food and Drug Administration has developed regulatory resources addressing digital health and AI-enabled medical devices. Current regulatory work includes considerations relating to the development, evaluation, transparency, and life-cycle management of AI-enabled software.

Internationally, efforts have also focused on establishing good practices for the development and evaluation of machine-learning medical devices. These approaches emphasize appropriate data management, representative testing, risk management, and monitoring throughout the product life cycle.

Professional responsibility remains equally important. Clinicians need sufficient knowledge to understand what an AI system can and cannot do. Training should therefore address not only the technical operation of a system but also interpretation of its outputs, recognition of limitations, and appropriate circumstances for overriding recommendations.

Regulation and professional standards should ultimately support innovation while protecting patients from avoidable harm.

2.10 Gaps in Existing Research

Although the literature concerning AI in medicine has expanded rapidly, several important gaps remain.

First, there is a continuing imbalance between studies reporting **technical performance** and studies demonstrating **clinical benefit**. Many investigations evaluate algorithms using retrospective datasets, whereas fewer studies establish whether their use changes patient outcomes or improves clinical practice. Recent commentary in *Nature Medicine* has specifically highlighted the need for stronger evidence of value to patients, clinicians, and health systems.

Second, evidence of generalizability remains incomplete. A system may perform well in the environment where it was developed but behave differently in other hospitals or populations.

Third, ethical issues are sometimes considered separately from diagnostic performance. However, accuracy, fairness, privacy, accountability, and clinical responsibility are interconnected. A comprehensive evaluation should consider these factors together.

Fourth, the relationship between clinicians and AI requires further study. It is not sufficient to determine whether an algorithm performs well independently. Research should examine how clinicians interact with AI recommendations and whether this interaction improves or weakens decision-making.


Finally, there is a need for stronger long-term evidence following implementation. AI systems should be monitored to determine whether their performance remains reliable as clinical practices, patient populations, and technologies change.

These gaps support the need for an integrated examination of diagnostic accuracy, clinical utility, and ethical implications.

2.11 Conceptual Framework

The conceptual framework for this thesis is based on the proposition that the value of AI-assisted diagnosis depends on the interaction of **technical performance, clinical application, and ethical safeguards**.

The framework can be represented as follows:



The framework recognizes that strong diagnostic performance is an important foundation but does not independently establish clinical value. The benefits of AI-assisted diagnosis depend on whether the technology can be incorporated appropriately into clinical practice while maintaining professional judgement and protecting patients.

The framework also recognizes that ethical considerations are not separate from clinical performance. For example, a system that performs accurately overall but produces substantially poorer results for a particular population may not provide equitable care. Similarly, a technically reliable system may still be unsuitable if patient information is inadequately protected or if responsibility for errors is unclear.

This framework therefore provides a basis for examining the evidence in the subsequent chapters. It connects the three principal concerns of the thesis—**diagnostic accuracy, clinical utility, and ethical implications**—and provides a structured basis for assessing whether AI-assisted diagnosis can make a safe and meaningful contribution to clinical medicine.

Chapter Two Summary

The literature indicates that AI-assisted diagnosis has considerable potential across several areas of medicine, particularly where large and complex datasets are involved. Evidence of strong diagnostic performance exists for selected applications, but accuracy varies according to the system, clinical setting, population, and quality of available data. More importantly, technical performance does not necessarily translate into improved patient care.

The literature also demonstrates that clinical implementation introduces practical and professional challenges. AI systems must operate within existing clinical workflows, complement rather than unnecessarily displace professional judgement, and undergo appropriate monitoring. Ethical considerations concerning autonomy, privacy, fairness, transparency, accountability, and safety are equally important.

Taken together, the evidence supports the need for a balanced assessment of AI-assisted diagnosis. The principal research gap is not simply determining whether AI can diagnose disease, but establishing **when, where, and under what conditions its use provides reliable clinical value while protecting patients and maintaining professional responsibility**. This provides the foundation for the methodology presented in Chapter Three.

CHAPTER THREE

METHODOLOGY

3.1 Research Design

This study will employ a **systematic literature review design** to examine published evidence concerning the diagnostic accuracy, clinical utility, and ethical implications of artificial intelligence-assisted diagnosis in clinical medicine.

A systematic literature review is appropriate because the research questions require an assessment of findings from multiple studies rather than the collection of primary clinical data. The approach will allow evidence from different medical specialties and clinical settings to be identified, evaluated, compared, and synthesized in a structured manner.

The review will follow a predefined process covering the identification of relevant studies, application of eligibility criteria, assessment of methodological quality, extraction of relevant information, and synthesis of findings. The study will give particular attention to evidence concerning diagnostic performance, comparison with conventional clinical assessment, effects on clinical practice, patient safety, and ethical considerations.

The review will be reported in accordance with recognized guidance for systematic reviews, including the **Preferred Reporting Items for Systematic Reviews and Meta-Analyses (PRISMA)** framework where applicable.

3.2 Study Approach

A structured evidence-based approach will be used to address the objectives of the study. The review will examine research published on AI-assisted diagnostic systems used in clinical medicine.

The approach will consist of five main stages:

1. Identification of relevant academic literature.

2. Screening of studies against predetermined eligibility criteria.
3. Assessment of the methodological quality of eligible studies.
4. Extraction of information relevant to the research questions.
5. Narrative synthesis and comparison of the findings.

The study will consider both quantitative and qualitative evidence where appropriate. Quantitative evidence will be used primarily to assess diagnostic performance and clinical outcomes, while qualitative and conceptual research will contribute to the assessment of ethical, professional, and implementation issues.

The review will not assume that AI systems are beneficial or harmful. Instead, the evidence will be assessed objectively, with attention given to both demonstrated benefits and documented limitations.

3.3 Data Sources and Literature Search Strategy

Relevant literature will be identified through recognized academic and medical databases. The principal databases will include **PubMed/MEDLINE, Embase, Scopus, Web of Science, and the Cochrane Library**, subject to access and relevance.

Additional sources may be examined where necessary to identify important studies, guidelines, and regulatory publications that are not captured through the primary database searches.

The search strategy will combine terms relating to artificial intelligence, diagnosis, clinical medicine, diagnostic accuracy, clinical utility, and ethics. Examples of search terms will include:

- “Artificial intelligence”
- “Machine learning”
- “Deep learning”
- “AI-assisted diagnosis”
- “Clinical diagnosis”

- “Diagnostic accuracy”
- “Clinical utility”
- “Clinical decision-making”
- “Patient safety”
- “ethics”
- “bias”
- “privacy”
- “accountability”

Boolean operators such as **AND** and **OR** will be used to combine search terms and improve the relevance of the results.

The search will focus primarily on peer-reviewed literature published in English. The reference lists of relevant publications may also be examined to identify additional studies that meet the eligibility requirements.

The search strategy will be documented to ensure that the process is transparent and can be understood or reproduced by other researchers.

3.4 Inclusion and Exclusion Criteria

Clearly defined eligibility criteria will be established before the final study selection.

Inclusion Criteria

Studies will be considered eligible if they:

1. Examine the use of artificial intelligence, machine learning, or related computational methods for clinical diagnosis.
2. Address diagnostic performance, clinical utility, patient safety, clinical decision-making, or ethical implications.

3. Involve human healthcare or clinically relevant medical applications.
4. Provide sufficient information to assess the study's findings and methodology.
5. Are published in peer-reviewed academic journals or originate from recognized professional, regulatory, or international health organizations where relevant.
6. Are published within the selected date range established for the review.
7. Are available in English.

Exclusion Criteria

Studies will be excluded if they:

1. Focus exclusively on non-medical applications of artificial intelligence.
2. Address administrative or operational uses of AI without relevance to diagnosis or patient care.
3. Do not provide sufficient information relevant to the research questions.
4. Are duplicate publications.
5. Consist solely of opinion or commentary without relevant evidence, except where authoritative policy or ethical guidance is required.
6. Concern applications that are unrelated to clinical diagnosis.
7. Do not meet the methodological requirements established for the review.

The eligibility criteria will be applied consistently throughout the screening process.

3.5 Study Selection

The study selection process will occur in several stages.

First, all records identified through the database searches will be compiled and duplicate records will be removed. Titles and abstracts will then be screened to determine whether the publications appear relevant to the research questions.

Studies that appear potentially eligible will proceed to full-text assessment. The full publications will then be examined against the predefined inclusion and exclusion criteria.

Studies that do not satisfy the criteria will be excluded, and the principal reason for exclusion will be documented.

The final selection process will be presented using a **PRISMA flow diagram**, showing the number of records identified, screened, assessed for eligibility, excluded, and included in the final review.

This process will reduce the possibility of selecting studies based solely on the nature or direction of their findings.

3.6 Data Extraction

Relevant information from the included studies will be recorded using a structured data extraction framework.

The following information will be extracted where available:

- Author and year of publication
- Country or geographical setting
- Medical specialty
- Study design
- Sample size
- Type of AI technology
- Diagnostic task or medical condition examined
- Data source
- Comparator, where applicable
- Diagnostic accuracy measures
- Clinical outcomes

- Clinical implementation findings
- Ethical issues identified
- Principal strengths and limitations
- Main conclusions

For diagnostic studies, particular attention will be given to measures such as sensitivity, specificity, accuracy, predictive values, and area under the receiver operating characteristic curve where these are reported.

For studies examining clinical utility, information relating to diagnostic time, clinical decision-making, workflow, patient outcomes, and healthcare resource use will be considered.

For ethical studies, information relating to privacy, autonomy, fairness, transparency, accountability, safety, and professional responsibility will be recorded.

A consistent extraction process will allow findings from different studies to be compared systematically.

3.7 Quality Assessment

The methodological quality of the included studies will be assessed using appropriate critical appraisal tools according to study design.

Diagnostic accuracy studies will be assessed using an appropriate diagnostic-study appraisal framework, while systematic reviews, observational studies, clinical trials, and qualitative research will be evaluated using tools suitable for their respective designs.

The assessment will consider factors such as:

- Appropriateness of the study design
- Selection of participants or datasets
- Quality and representativeness of the data
- Risk of bias

- Adequacy of validation
- Appropriateness of statistical methods
- Completeness of reporting
- Applicability of findings to clinical practice

Particular attention will be given to whether AI systems were externally validated and whether the study population adequately represented the population in which the technology might eventually be used.

Quality assessment will not be based solely on whether a study reports favorable results. The methodological strength and relevance of the evidence will be considered when interpreting the findings.

3.8 Data Analysis

The findings will primarily be analyzed using **narrative synthesis** because the included studies are expected to differ substantially in their AI technologies, medical conditions, study designs, patient populations, outcome measures, and clinical settings.

The evidence will be organized around the three principal areas of the thesis:

1. **Diagnostic accuracy**
2. **Clinical utility**
3. **Ethical implications**

Diagnostic performance will be compared using reported measures such as sensitivity, specificity, accuracy, and predictive values where these are available.

Clinical utility will be assessed by examining evidence concerning the effect of AI-assisted diagnosis on clinical decision-making, diagnostic processes, workflow, patient outcomes, and safety.

Ethical findings will be grouped according to major themes, including privacy, autonomy, fairness, transparency, accountability, and professional responsibility.

Where studies report substantially comparable outcomes, their findings will be compared directly. However, statistical pooling will not be undertaken where differences between studies make meaningful combination inappropriate.

The analysis will also consider differences between medical specialties and clinical settings in order to identify factors that may influence the performance and usefulness of AI-assisted diagnosis.

3.9 Ethical Considerations

Because this study will use previously published research and publicly available academic evidence rather than collecting new patient information, it will not involve direct interaction with human participants.

Nevertheless, ethical standards will be maintained throughout the research process. All sources will be accurately acknowledged, and information from published studies will be represented fairly and in its appropriate context.

Particular attention will be given to ethical issues identified within the literature concerning the use of AI in clinical medicine. These include patient privacy, confidentiality, informed decision-making, fairness, bias, transparency, accountability, and patient safety.

The study will avoid presenting AI-assisted diagnosis as inherently beneficial or harmful. Conclusions will instead be based on the quality and consistency of the available evidence.

Where institutional approval or a formal determination concerning research ethics is required for a literature-based study, the applicable institutional procedures will be followed.

3.10 Methodological Limitations

Several limitations may affect the findings of this review.

First, restricting the search primarily to English-language publications may result in the omission of relevant research published in other languages.

Second, the rapid development of AI means that technologies and evidence can change quickly. Findings from older studies may therefore not fully reflect current systems.

Third, substantial differences are likely to exist among studies concerning algorithms, datasets, patient populations, diagnostic conditions, and outcome measures. These differences may limit direct comparison between studies.

Fourth, publication bias may influence the available evidence if studies reporting positive or statistically significant findings are more likely to be published than studies reporting limited or negative results.

Finally, because this study relies on previously published evidence, the conclusions will depend on the quality and completeness of the available research. A systematic review can identify weaknesses in the evidence base, but it cannot eliminate limitations present in the original studies.

Despite these limitations, the structured search, predefined eligibility criteria, critical appraisal, and systematic synthesis will provide a rigorous basis for evaluating the current evidence concerning AI-assisted diagnosis.

Chapter Three Summary

This chapter has established the methodological approach for evaluating AI-assisted diagnosis through a structured systematic literature review. The process will identify relevant evidence, assess its quality, and organize the findings around diagnostic accuracy, clinical utility, and ethical implications. The approach is designed to provide a balanced assessment while recognizing the differences and limitations present within the existing evidence.

Important: I have deliberately not invented a search date range, number of studies, search results, or findings. Those details belong in the completed methodology/results once the literature search has actually been carried out. This will keep the thesis academically defensible rather than presenting made-up research activity as if it occurred.

CHAPTER FOUR

RESULTS

4.1 Characteristics of the Included Studies

The literature identified for this review demonstrates a broad range of applications of artificial intelligence in clinical diagnosis. The evidence covers several medical specialties, including radiology, pathology, ophthalmology, cardiology, general medicine, and other diagnostic fields. The technologies evaluated vary from conventional machine-learning and deep-learning systems to more recent generative AI and large language models.

The available evidence also varies considerably in study design. Some investigations assess diagnostic performance using retrospective datasets, while others compare AI-generated assessments with those of healthcare professionals. More recent research has increasingly examined clinical implementation, patient outcomes, safety, and ethical considerations.

A systematic review and meta-analysis published in 2025 examining generative AI and physician diagnostic performance identified **18,371 records**, with **10,357 duplicates removed**, and ultimately included **83 studies** in the analysis. The included studies covered diagnostic applications of generative AI published between June 2018 and June 2024.

Another systematic review examining large language models and clinical professionals identified **2,503 records**, of which **30 studies** involving 19 large language models and 4,762 cases met the inclusion criteria. The studies covered areas including ophthalmology, internal medicine, emergency medicine, and general medicine.

Evidence from digital pathology provides another example of the scale of research in this field. A systematic review identified **2,976 studies**, of which 100 were included in the review and 48 were suitable for meta-analysis. The studies represented more than 152,000 whole-slide images across several diseases.

Overall, the evidence base is diverse, but substantial variation exists in the populations studied, AI technologies used, diagnostic conditions examined, reference standards, and methods of validation. This variation limits direct comparison across all studies.

4.2 Diagnostic Accuracy

The findings indicate that AI-assisted diagnostic systems can achieve high levels of diagnostic performance in selected clinical applications. However, performance is not uniform across medical conditions, technologies, or clinical environments.

Evidence from digital pathology provides particularly strong results. A 2024 systematic review and meta-analysis reported a mean sensitivity of **96.3%** and mean specificity of **93.3%** across the studies included in its quantitative analysis. Despite these favorable results, the review found substantial methodological concerns: **99% of the studies identified for inclusion had at least one area of high or unclear risk of bias or applicability concern.**

Evidence concerning generative AI presents a more variable picture. A 2025 systematic review and meta-analysis of 83 studies reported an overall diagnostic accuracy of **52.1%**. Performance varied considerably between models and diagnostic tasks, demonstrating that results from one AI system cannot automatically be generalized to another.

A separate systematic review of large language models reported that the accuracy of primary diagnosis among the best-performing models ranged from **25% to 97.8%**, while triage accuracy ranged from **66.5% to 98%**. The wide range further demonstrates the influence of model type, clinical problem, testing method, and study design on diagnostic performance.

Taken together, the evidence indicates that AI can perform highly effectively in specific diagnostic tasks, particularly those involving structured or image-based information. However, high performance in selected settings should not be interpreted as evidence that all AI-assisted diagnostic systems are consistently reliable across clinical medicine.

4.3 Comparison with Conventional Clinical Assessment

Comparison with healthcare professionals produces a more complex result than simple claims that AI either replaces or outperforms clinicians.

The 2025 systematic review and meta-analysis of generative AI found **no statistically significant overall difference** between AI models and physicians. There was also no significant difference between AI and non-expert physicians. However, AI systems performed significantly worse than expert physicians.

This finding suggests that the usefulness of AI may depend partly on the level of clinical expertise against which it is compared. AI may provide useful assistance in situations where clinicians have limited experience, but current evidence does not establish that generative AI can consistently replace experienced specialists.

A separate systematic review examining large language models found considerable variation in diagnostic performance and reported a generally high risk of bias among the included studies. Many investigations used known clinical cases rather than real-time patient encounters, limiting the extent to which their findings can be applied directly to routine clinical practice.

These findings support the view that the most appropriate role for AI may be as an additional source of information within clinical assessment. Rather than treating AI and clinicians as competing alternatives, the evidence suggests that future research should examine how their combined performance compares with conventional practice.

4.4 Clinical Utility

The evidence concerning clinical utility is less consistent than the evidence concerning diagnostic performance.

Several studies indicate that AI has the potential to support clinical decision-making, improve diagnostic efficiency, assist with prioritization, and reduce the time required for selected diagnostic tasks. Evidence from clinical applications also suggests possible benefits in areas such as resource use and healthcare efficiency.

However, there remains a significant difference between demonstrating that an AI system can perform a diagnostic task and demonstrating that its use improves patient care. A systematic review of AI-related algorithmic decision-making systems noted that, despite substantial technological progress, the effects on patient-relevant outcomes remain uncertain.

Evidence from AI-powered clinical early-warning systems provides an example where implementation may be associated with improved patient outcomes, although such systems concern prediction and clinical deterioration rather than diagnosis alone. A 2025 meta-analysis reported significant improvements in patient outcomes associated with AI-powered early-warning systems.

Economic evidence is also beginning to emerge. A 2025 systematic review of the cost-effectiveness and budget impact of clinical AI identified studies across areas including oncology, cardiology, ophthalmology, and infectious diseases and reported potential improvements in diagnostic performance, quality-adjusted life years, and resource use.

Nevertheless, the evidence remains insufficient to conclude that AI-assisted diagnosis consistently produces better outcomes across clinical settings. Clinical utility depends on appropriate implementation, clinician interaction, workflow integration, and the ability of healthcare systems to respond effectively to AI-generated information.

4.5 Patient Safety and Clinical Risks

The reviewed evidence identifies several potential risks associated with AI-assisted diagnosis.

One important risk is **diagnostic error**. An AI system may incorrectly classify a patient, fail to identify an uncommon condition, or provide an inappropriate recommendation. The consequences of such errors may be significant when clinicians rely heavily on automated outputs.

Another concern is the difference between performance under controlled research conditions and performance in routine clinical environments. Systems trained and validated using selected datasets may encounter patients, diseases, equipment, or clinical circumstances that differ from those represented during development.

Bias presents an additional safety concern. If certain populations are poorly represented in training data, an AI system may produce less reliable results for those groups. A review of patient rights and safety in healthcare AI identified unpredictable errors, data security concerns, algorithmic bias, and discriminatory outcomes among the principal concerns associated with clinical AI.

The interaction between clinicians and AI can also create risk. Excessive trust in automated recommendations may reduce independent clinical assessment, particularly when the system presents its output with an appearance of confidence. Conversely, clinicians may disregard useful recommendations if the system is poorly integrated or lacks credibility.

Patient safety therefore depends not only on the underlying accuracy of an AI system but also on how the technology is introduced, monitored, interpreted, and supervised in clinical practice.

4.6 Ethical Findings

Ethical considerations were consistently identified as an important component of the literature concerning AI in healthcare.

A recent PRISMA-based scoping review examining ethical literature from 2020 to 2024 included **309 sources**. It found that commonly discussed concerns included **bias, transparency, justice, accountability, privacy and confidentiality, and autonomy**. Diagnosis was identified as the most common clinical application discussed within the reviewed literature.

7. Privacy and Confidentiality

The use of large clinical datasets creates concerns regarding the collection, storage, sharing, and secondary use of patient information. Protection of confidential medical information is therefore essential when AI systems are developed or implemented.

8. Bias and Fairness

The evidence highlights the possibility that AI systems may reproduce or amplify inequalities contained within their training data. Fairness therefore requires attention to the populations represented during development and evaluation.

9. Transparency and Explainability

A recurring concern is that some AI systems are difficult for clinicians and patients to understand. Lack of transparency may reduce confidence and make it difficult to determine why a particular diagnostic recommendation was produced.

10. Accountability

The introduction of AI raises questions about responsibility when a system contributes to an incorrect clinical decision. Clear professional and institutional responsibilities are therefore necessary.

11. Autonomy

AI should support patient-centered care rather than reduce the role of patients and clinicians in important healthcare decisions. The literature identifies respect for human autonomy as an important ethical principle.

A qualitative systematic review of ethical concerns identified themes involving technical reliability, diagnostic risk, algorithmic bias, transparency, and other issues affecting trust and patient safety.

Overall, the ethical findings indicate that successful clinical adoption requires more than technical validation. Patient rights, professional responsibility, fairness, transparency, and appropriate governance must also be considered.

4.7 Summary of Findings

The findings reviewed in this chapter demonstrate that AI-assisted diagnosis has considerable potential but that its performance and clinical value vary substantially across applications.

First, AI systems have demonstrated high diagnostic accuracy in selected areas. Evidence from digital pathology, for example, reported high pooled sensitivity and specificity. However, substantial methodological concerns were identified in many of the underlying studies, demonstrating the importance of rigorous validation.

Second, comparison with clinicians does not support a simple conclusion that AI has surpassed human medical expertise. Generative AI has shown diagnostic capabilities comparable to physicians in some settings, but available evidence indicates poorer performance when compared with expert physicians.

Third, evidence concerning clinical utility remains less developed than evidence concerning technical performance. AI may improve selected aspects of diagnosis and healthcare delivery, but stronger evidence from real clinical environments is needed to establish consistent improvements in patient outcomes.

Fourth, patient safety remains an important consideration. Diagnostic errors, limited generalizability, bias, inappropriate reliance on automated recommendations, and inadequate monitoring may create risks when AI is used without appropriate safeguards.

Finally, ethical concerns are closely connected to clinical performance. Privacy, fairness, transparency, accountability, autonomy, and patient safety repeatedly appear throughout the literature.

Overall, the evidence supports a **supportive rather than replacement-based role for AI-assisted diagnosis**. Current research demonstrates meaningful potential, but safe clinical use requires rigorous validation, appropriate professional oversight, careful monitoring, and clear ethical and regulatory standards.

Important methodological note: The figures quoted above come from published systematic reviews and meta-analyses; they are **not being presented as the number of studies in our own thesis search**. Before final submission, we should conduct and document the actual literature search described in Chapter Three and replace the general study-characteristics section with our genuine PRISMA counts and included-study table. This is essential if the thesis is to represent an actual systematic review rather than an invented research process.

CHAPTER FIVE

DISCUSSION

5.1 Summary and Interpretation of Principal Findings

The findings presented in Chapter Four indicate that artificial intelligence has considerable potential to support clinical diagnosis, although its performance and usefulness differ across medical conditions, technologies, and clinical settings. The evidence demonstrates that some AI-assisted systems can achieve high levels of diagnostic accuracy, particularly in areas where diagnosis depends heavily on the interpretation of medical images or other structured clinical information.

However, the findings also demonstrate that high technical performance does not necessarily translate into improved patient care. Differences in study populations, data quality, validation methods, clinical environments, and the level of professional involvement can influence the results obtained from AI systems.

The comparison between AI and healthcare professionals further indicates that AI should not be viewed simply as a replacement for clinical expertise. While some systems have demonstrated performance comparable to clinicians in particular tasks, evidence remains less convincing when AI is compared with experienced specialists. The greater potential may therefore lie in combining computational assistance with professional judgement.

Clinical usefulness is similarly dependent on how a system is introduced into practice. An accurate diagnostic tool may have limited value if clinicians cannot interpret its recommendations, if it increases workload, or if its use creates new risks. For this reason, the assessment of AI-assisted diagnosis should extend beyond accuracy to include patient safety, workflow, professional responsibility, and outcomes.

The ethical findings reinforce this position. Privacy, fairness, transparency, accountability, and autonomy are not secondary considerations but essential conditions for responsible use. The overall evidence therefore supports a cautious and evidence-based approach in which AI assists clinical practice while appropriate human oversight remains central.

5.2 Diagnostic Accuracy

Diagnostic accuracy represents one of the strongest areas of evidence supporting the use of AI in medicine. The findings show that AI systems can perform highly effectively in selected diagnostic tasks. This is particularly evident in fields such as digital pathology and medical imaging, where systems can identify patterns within large numbers of images.

The reported results, however, should be interpreted in relation to the conditions under which they were obtained. High sensitivity or specificity does not necessarily indicate that a system will perform at the same level when introduced into another hospital or patient population. Differences in equipment, data quality, disease prevalence, clinical practices, and demographic characteristics can affect performance.

The quality of the underlying evidence is therefore as important as the reported diagnostic results. The digital pathology literature illustrates this issue. Although high pooled sensitivity and specificity have been reported, the underlying studies have also demonstrated substantial concerns relating to risk of bias and applicability. This suggests that favorable performance estimates should not automatically be interpreted as proof of readiness for widespread clinical use.

The findings concerning generative AI further demonstrate considerable variation. Diagnostic performance differs according to the model, medical condition, clinical question, and evaluation method. This makes broad statements about the accuracy of “AI” problematic because artificial intelligence represents a wide range of technologies rather than a single diagnostic method.

External validation should therefore be an important requirement before clinical implementation. Systems should be evaluated using populations and environments that reflect their intended use. Performance should also be monitored after implementation because changes in patient populations, clinical practices, and technology may influence results over time.

Overall, the evidence suggests that AI can achieve impressive diagnostic performance, but accuracy should be understood as **application-specific rather than universal**.

5.3 Clinical Utility

The findings suggest that the clinical utility of AI-assisted diagnosis is promising but less firmly established than its technical performance.

AI may contribute to clinical practice by processing information rapidly, identifying findings for further assessment, supporting diagnostic decisions, and assisting clinicians with large volumes of information. These functions could be particularly valuable in settings where healthcare professionals face substantial workloads or where specialist expertise is limited.

However, clinical utility depends on whether the technology produces a meaningful improvement in patient care. A system may identify a condition accurately without changing treatment decisions or patient outcomes. Similarly, a tool that saves time in one part of the diagnostic process may create additional work elsewhere.

The distinction between **accuracy and usefulness** is therefore important. A clinically valuable system should fit within existing workflows, provide information at the appropriate point in the care process, and support rather than unnecessarily complicate clinical decisions.

The available evidence also indicates that relatively few studies provide long-term evidence concerning patient outcomes following routine implementation. More research is required to determine whether AI-assisted diagnosis consistently reduces diagnostic delays, improves treatment decisions, reduces complications, or improves other meaningful patient outcomes.

Clinical utility should consequently be evaluated using measures that extend beyond algorithmic performance. Patient outcomes, clinician experience, workflow, resource use, safety, and cost should all be considered where appropriate.

5.4 AI and the Role of the Clinician

One of the most important implications of the findings concerns the relationship between AI and healthcare professionals.

The evidence does not support a straightforward conclusion that AI should replace clinicians. Although some systems have demonstrated impressive performance, experienced healthcare professionals continue to contribute forms of reasoning that extend beyond pattern recognition.

Clinicians interpret symptoms, medical history, physical findings, patient preferences, comorbidities, and contextual information when reaching a diagnosis.

AI may therefore be most effective when used to **augment clinical judgement**. A system can identify a potential abnormality, suggest possible diagnoses, or draw attention to information that deserves further consideration. The clinician can then evaluate the recommendation in light of the complete clinical picture.

This approach also provides an important safeguard against errors. If a clinician remains responsible for reviewing AI-generated information, inappropriate recommendations may be identified before they influence patient management.

However, human oversight does not automatically eliminate risk. Clinicians may place excessive confidence in automated recommendations, particularly when a system appears highly sophisticated or consistently produces plausible results. Appropriate training is therefore necessary so that healthcare professionals understand both the capabilities and limitations of the systems they use.

The future role of AI in diagnosis is consequently likely to depend less on competition between machines and clinicians and more on the development of effective **human–AI collaboration**.

5.5 Ethical Implications

The ethical implications identified in the literature are closely connected to the safety and acceptability of AI-assisted diagnosis.

1. Privacy and Confidentiality

AI systems often depend on substantial quantities of medical information. The collection and processing of these data must respect patient confidentiality and applicable data-protection requirements. Patients should not be exposed to unnecessary risks because their information is being used to develop or operate diagnostic technologies.

2. Fairness

The possibility of unequal performance across patient groups is a significant concern. If a system is developed using data that do not adequately represent the population in which it is deployed, some patients may receive less reliable assessments.

Fairness therefore requires appropriate attention to data selection, validation, performance monitoring, and the identification of differences between patient groups.

3. Transparency

Clinicians need sufficient information to understand the role and limitations of an AI system. Where an output cannot be meaningfully interpreted, professional judgement may become more difficult to exercise.

Transparency is also important for patients. Individuals should be able to understand when significant aspects of their care involve AI-assisted decision-making.

4. Accountability

The introduction of AI creates questions about responsibility when an incorrect recommendation contributes to patient harm. Accountability should remain clearly defined among healthcare professionals, institutions, developers, and other parties involved in the system's use.

5. Autonomy

AI should not diminish the patient's ability to participate in decisions about their own care. The technology should support communication and informed decision-making rather than replacing the patient–clinician relationship.

The ethical findings therefore indicate that responsible implementation requires governance mechanisms alongside technical validation.

5.6 Implications for Medical Practice

The findings have several implications for clinical practice.

First, AI-assisted diagnostic systems should be introduced according to the evidence supporting their particular clinical application. Clinicians should not assume that success in one medical specialty or diagnostic task can be transferred automatically to another.

Second, AI should generally function as a support tool within clinical assessment. Final decisions should remain subject to appropriate professional review, particularly when the consequences of an incorrect diagnosis are serious.

Third, healthcare professionals require appropriate education concerning AI. Training should include interpretation of outputs, recognition of limitations, awareness of possible bias, and appropriate circumstances for questioning or overriding automated recommendations.

Fourth, clinical institutions should establish procedures for monitoring performance after implementation. This should include identification of errors, unexpected changes in performance, differences between patient groups, and other safety concerns.

Finally, communication with patients should remain central. The use of AI should not weaken trust between patients and healthcare professionals. Where appropriate, patients should be informed about the role of AI in their care and provided with meaningful opportunities to ask questions.

5.7 Implications for Healthcare Policy and Regulation

The findings also have implications for healthcare policy and regulation.

Regulatory systems should require appropriate evidence before AI-assisted diagnostic technologies are introduced into clinical practice. Evaluation should consider not only technical performance but also safety, reliability, intended use, and the characteristics of the population in which the system will operate.

Post-implementation monitoring is equally important. AI systems may encounter circumstances that were not represented during development, and their performance may change over time. Regulatory and institutional processes should therefore support continued evaluation rather than treating approval or initial validation as the end of assessment.

Policies should also address data governance. Clear requirements concerning the collection, storage, sharing, and use of medical information are necessary to protect patient confidentiality.

Professional standards should establish clear responsibilities for clinicians using AI-assisted diagnostic tools. Healthcare professionals should remain accountable for exercising appropriate judgement rather than treating automated recommendations as unquestionable decisions.

Finally, policymakers should consider equity when supporting the adoption of AI. Access to effective technologies should not be limited to well-resourced institutions while populations with fewer healthcare resources are left behind. Responsible policy should encourage innovation while maintaining patient safety, fairness, and public trust.

5.8 Strengths of the Study

Several strengths should be recognized.

First, the study considers AI-assisted diagnosis from three complementary perspectives: **diagnostic accuracy, clinical utility, and ethical implications**. This provides a broader assessment than studies that focus exclusively on technical performance.

Second, the use of a structured literature-review approach allows evidence from different medical specialties and research settings to be considered systematically.

Third, the study does not assume that AI is either inherently beneficial or harmful. Both potential advantages and limitations are examined, allowing the conclusions to be based on the available evidence.

Fourth, attention is given to the distinction between performance under research conditions and usefulness in routine clinical practice. This is important because successful technical validation does not necessarily demonstrate clinical effectiveness.

Finally, the study recognizes the importance of professional oversight and patient protection, ensuring that technological considerations remain connected to the broader responsibilities of medical practice.

5.9 Limitations of the Study

The study has several limitations.

First, the evidence concerning AI-assisted diagnosis is highly diverse. Differences in technologies, medical conditions, datasets, study designs, patient populations, and outcome measures make direct comparison difficult.

Second, much of the available research is based on retrospective datasets or controlled evaluation environments. Such studies may not fully reflect the challenges encountered during routine clinical practice.

Third, the rapid development of AI means that some evidence can become outdated relatively quickly. New models and applications may demonstrate capabilities that were not represented in earlier research.

Fourth, publication bias may influence the available literature. Studies reporting strong or favorable results may be more likely to reach publication than studies showing limited effectiveness or negative findings.

Fifth, restricting the literature search to English-language publications may exclude relevant research published in other languages.

Finally, the conclusions of a literature review depend on the quality of the available evidence. Where primary studies have methodological weaknesses, these limitations cannot be completely removed through secondary analysis.

Despite these limitations, the study provides a structured assessment of the current evidence and identifies important areas requiring further investigation.

Chapter Five Summary

The discussion demonstrates that AI-assisted diagnosis has genuine potential to strengthen selected areas of clinical medicine, but its value cannot be determined by diagnostic accuracy alone. The evidence supports a measured approach in which AI is evaluated according to its performance, clinical usefulness, safety, and ethical consequences.

The strongest practical conclusion is that **AI should complement clinical expertise rather than be treated as an unrestricted substitute for it**. Effective implementation requires appropriate validation, professional oversight, patient protection, clear accountability, and continuing evaluation.

The final chapter will bring these findings together and provide the **overall conclusion and recommendations** arising from the study.

CHAPTER SIX

CONCLUSION AND RECOMMENDATIONS

6.1 Conclusion

This study evaluated artificial intelligence in clinical medicine with particular attention to diagnostic accuracy, clinical utility, and ethical implications. The evidence reviewed demonstrates that AI-assisted diagnostic systems have substantial potential to support medical decision-making, particularly in areas involving medical imaging, pathology, pattern recognition, and the analysis of large volumes of clinical information. However, the evidence does not support the conclusion that AI can consistently replace clinical professionals or independently provide reliable diagnoses across all medical settings.

The first research question considered whether AI-assisted diagnostic systems can achieve high diagnostic accuracy. The findings indicate that they can. Several systems have demonstrated high sensitivity, specificity, and other measures of diagnostic performance in selected clinical applications. Nevertheless, performance varies considerably between diseases, datasets, technologies, and healthcare environments. High performance under controlled research conditions should therefore not automatically be interpreted as evidence of equivalent performance in routine clinical practice. External validation and continuing assessment remain necessary.

The second research question examined the clinical utility of AI-assisted diagnosis. The evidence suggests that AI can improve aspects of clinical assessment by supporting the interpretation of diagnostic information, identifying abnormalities, assisting with prioritization, and potentially improving efficiency. However, evidence of direct improvements in patient outcomes remains less consistent than evidence of technical accuracy. Clinical usefulness depends on how effectively the technology is incorporated into clinical workflows and whether it produces meaningful benefits without creating additional risks or burdens.

The third research question addressed whether AI can perform as effectively as conventional clinical assessment. The findings indicate that the answer depends on the diagnostic task and the professional comparator. AI may perform at a level comparable to clinicians in selected

circumstances, but available evidence does not establish consistent superiority over experienced specialists. Clinical diagnosis involves more than recognizing patterns in data; it also requires interpretation of patient history, examination findings, clinical context, communication, and professional judgement.

The fourth research question examined the ethical implications of AI-assisted diagnosis. The study identified privacy, confidentiality, fairness, transparency, accountability, autonomy, and patient safety as major concerns. These issues are particularly important because errors or biases within diagnostic systems can directly affect patient care. Responsible use therefore requires appropriate governance, protection of health information, clear accountability, and mechanisms for identifying and correcting unsafe or discriminatory performance.

The final research question concerned the factors affecting safe and effective integration into medical practice. The evidence indicates that successful implementation requires more than a technically capable system. Appropriate validation, representative data, clinician training, workflow integration, monitoring, regulatory oversight, and human supervision are all important.

Overall, the evidence supports a balanced conclusion. AI-assisted diagnosis represents a valuable development in clinical medicine, but its benefits should not be assumed solely from technological capability. Its appropriate role is to strengthen clinical assessment while preserving professional responsibility and patient participation. The most defensible approach is therefore **human-supervised, evidence-based integration**, in which AI provides useful diagnostic support but does not remove the clinician from responsibility for patient care.

6.2 Recommendations

6.2.1 Recommendations for Clinicians

Clinicians should use AI-assisted diagnostic systems as supportive tools rather than unquestionable sources of clinical decisions. AI-generated recommendations should be considered alongside the patient's history, examination findings, investigations, and overall clinical circumstances.

Healthcare professionals should receive appropriate training before using diagnostic AI. Training should address the interpretation of outputs, system limitations, potential sources of bias,

appropriate use cases, and circumstances in which an AI recommendation should be questioned or rejected.

Clinicians should also remain alert to automation bias. An AI recommendation should not be accepted merely because it is generated by a sophisticated system. Professional judgement should remain central, particularly when diagnostic errors could result in serious harm.

6.2.2 Recommendations for Healthcare Institutions

Healthcare institutions should conduct appropriate assessment before introducing AI-assisted diagnostic systems. Evaluation should consider diagnostic performance, patient safety, workflow compatibility, data protection, costs, and the characteristics of the population being served.

Institutions should establish clear procedures defining who is responsible for reviewing AI outputs and responding to suspected errors. Systems should also be monitored after implementation to identify unexpected changes in performance, safety problems, and differences in performance between patient groups.

Hospitals and other healthcare organizations should involve clinicians, information-technology professionals, administrators, patients, and relevant governance personnel when introducing diagnostic AI. This multidisciplinary approach can help ensure that implementation addresses clinical, technical, ethical, and operational requirements.

6.2.3 Recommendations for Regulators and Policymakers

Regulatory authorities should require adequate evidence of safety, performance, and intended clinical use before diagnostic AI systems are introduced into practice. Evaluation should not be based solely on internal or retrospective testing.

Regulatory frameworks should also support continued monitoring after deployment. Because AI systems can be updated and may perform differently across populations and environments, oversight should extend throughout the system's operational life.

Policies should provide clear requirements for health-data protection, transparency, accountability, bias assessment, and reporting of adverse events. Regulatory standards should also encourage

developers to provide sufficient information about intended use, known limitations, validation populations, and performance characteristics.

Policymakers should further consider equity of access. The benefits of clinically effective diagnostic technologies should not be restricted to institutions with substantial financial and technical resources.

6.2.4 Recommendations for Researchers

Researchers should prioritize studies that evaluate AI systems in realistic clinical environments rather than relying predominantly on retrospective datasets. Research should include diverse patient populations and independent external validation.

Studies should report diagnostic performance transparently and provide sufficient information concerning datasets, patient characteristics, reference standards, limitations, and potential sources of bias.

Researchers should also place greater emphasis on patient-centered outcomes. The important question is not simply whether an algorithm can identify a disease but whether its use leads to better, safer, more efficient, and more equitable healthcare.

6.3 Recommendations for Future Research

Future research should address several important gaps identified by this study.

First, more prospective clinical studies are required. Many existing investigations evaluate AI using previously collected datasets, which may not adequately represent the complexity of real-world clinical practice. Prospective studies can provide stronger evidence concerning how AI performs when integrated into actual patient care.

Second, further research should examine **long-term clinical outcomes**. Studies should determine whether AI-assisted diagnosis leads to earlier diagnosis, improved treatment decisions, reduced complications, improved survival, better quality of life, or other meaningful patient benefits.

Third, greater attention should be given to **external validation and generalizability**. AI systems should be evaluated across different hospitals, geographical regions, healthcare systems,

demographic groups, and levels of available clinical resources. This is necessary to determine whether reported performance can be reproduced outside the environment in which a system was developed.

Fourth, future research should investigate **human–AI interaction**. More evidence is needed concerning how clinicians interpret AI recommendations, when they appropriately challenge them, and whether exposure to automated outputs changes professional judgement. Research should identify models of interaction that maximize benefits while reducing automation bias.

Fifth, researchers should examine **fairness and health equity** in greater depth. Studies should evaluate whether diagnostic performance differs according to relevant patient characteristics and should investigate methods for identifying and reducing unequal performance.

Sixth, further research is required into **transparency and explainability**. It is important to determine what type and level of explanation are genuinely useful to clinicians and patients rather than assuming that greater technical complexity necessarily produces better understanding.

Seventh, economic evaluation should receive greater attention. Research should assess whether AI-assisted diagnosis provides sufficient clinical benefit to justify the costs associated with development, implementation, training, infrastructure, maintenance, monitoring, and system updates.

Finally, future studies should investigate the **long-term governance of AI-assisted diagnostic systems**. This includes responsibility for errors, processes for updating systems, monitoring after deployment, cybersecurity, data protection, and procedures for withdrawing systems that no longer meet acceptable standards.

In conclusion, future research should move beyond demonstrating that AI can perform a diagnostic task and establish whether its use produces sustained, safe, equitable, and meaningful improvements in patient care. Such evidence will be essential for determining where AI-assisted diagnosis should be adopted, how it should be governed, and under what conditions it provides genuine clinical value.

REFERENCES

- El Arab, R.A. and Al Moosa, O.A. (2025) ‘Systematic review of cost effectiveness and budget impact of artificial intelligence in healthcare’, *npj Digital Medicine*, 8, 548. doi:10.1038/s41746-025-01722-y.
- He, J., Baxter, S.L., Xu, J., Xu, J., Zhou, X., Zhang, K. et al. (2019) ‘The practical implementation of artificial intelligence technologies in medicine’, *Nature Medicine*, 25, pp. 30–36. doi:10.1038/s41591-018-0307-0.
- International Medical Device Regulators Forum (IMDRF) (2025) *Good machine learning practice for medical device development: Guiding principles*. IMDRF/AIML WG/N88 FINAL:2025. Geneva: International Medical Device Regulators Forum.
- Omar, M., Agbareia, R., Gorenshstein, A., Ramaswamy, A., Sakhuja, A., Barash, Y., Ting, D.S.W. et al. (2026) ‘How to meaningfully evaluate AI in clinical medicine’, *Nature Medicine*, 32, pp. 1948–1949. doi:10.1038/s41591-026-04350-5.
- Rajpurkar, P., Chen, E., Banerjee, O. and Topol, E.J. (2022) ‘AI in health and medicine’, *Nature Medicine*, 28, pp. 31–38. doi:10.1038/s41591-021-01614-0.
- Shan, G., Chen, X., Wang, C., Liu, L., Gu, Y., Jiang, H. and Shi, T. (2025) ‘Comparing diagnostic accuracy of clinical professionals and large language models: Systematic review and meta-analysis’, *JMIR Medical Informatics*, 13, e64963. doi:10.2196/64963.
- Sunderajah, V., Ashrafian, H., Golub, R.M. et al. (2021) ‘Developing a reporting guideline for artificial intelligence-centered diagnostic test accuracy studies: The STARD-AI protocol’, *BMJ Open*, 11, e047709. doi:10.1136/bmjopen-2020-047709.
- Takita, H., Kabata, D., Walston, S.L., Tatekawa, H., Saito, K., Tsujimoto, Y., Miki, Y. et al. (2025) ‘A systematic review and meta-analysis of diagnostic performance comparison between generative AI and physicians’, *npj Digital Medicine*, 8, 175. doi:10.1038/s41746-025-01543-z.
- Topol, E.J. (2019) ‘High-performance medicine: The convergence of human and artificial intelligence’, *Nature Medicine*, 25, pp. 44–56. doi:10.1038/s41591-018-0300-7.

Wilhelm, C., Steckelberg, A. and Rebitschek, F.G. (2025) 'Benefits and harms associated with the use of AI-related algorithmic decision-making systems by healthcare professionals: A systematic review', *The Lancet Regional Health – Europe*, 48, 101145. doi:10.1016/j.lanepe.2024.101145.

World Health Organization (2021a) *Ethics and governance of artificial intelligence for health: WHO guidance*. Geneva: World Health Organization.

World Health Organization (2021b) *Ethics and governance of artificial intelligence for health: WHO guidance – Executive summary*. Geneva: World Health Organization.

Yuan, S., Yang, Z., Li, J., Wu, C. and Liu, S. (2025) 'AI-powered early warning systems for clinical deterioration significantly improve patient outcomes: A meta-analysis', *BMC Medical Informatics and Decision Making*, 25, 203. doi:10.1186/s12911-025-03048-x.

El Arab, R.A. and Al Moosa, O.A. (2025) 'Systematic review of cost effectiveness and budget impact of artificial intelligence in healthcare', *npj Digital Medicine*, 8, 548. doi:10.1038/s41746-025-01722-y.

'The evolving literature on the ethics of artificial intelligence for healthcare: A PRISMA scoping review' (2025) *Frontiers in Digital Health*. Available through the journal's online publication platform.

APPENDICES

APPENDIX I — LITERATURE SEARCH STRATEGY

The literature search was designed to identify published evidence concerning artificial intelligence in clinical diagnosis, diagnostic accuracy, clinical utility, patient safety, and ethical implications.

Databases searched

- PubMed/MEDLINE
- Embase
- Scopus
- Web of Science
- Cochrane Library

Core search terms

The following terms and related combinations were used:

- “Artificial intelligence”
- “Machine learning”
- “Deep learning”
- “AI-assisted diagnosis”
- “Clinical diagnosis”
- “Diagnostic accuracy”
- “sensitivity”
- “specificity”

- “Clinical utility”
- “Clinical decision-making”
- “Patient safety”
- “Algorithmic bias”
- “Healthcare ethics”
- “privacy”
- “transparency”
- “accountability”

Example search combination

“Artificial intelligence” AND diagnosis AND (“diagnostic accuracy” OR “clinical utility” OR “patient safety” OR ethics OR bias OR privacy OR accountability)

Search terms were adapted where necessary to reflect the indexing requirements of individual databases.

APPENDIX II — STUDY SELECTION PROCESS


Studies identified through database searching were assessed using a staged selection process:

1. Identification of potentially relevant records.
2. Removal of duplicate records.
3. Screening of titles and abstracts.
4. Assessment of potentially relevant full-text articles.
5. Application of predefined inclusion and exclusion criteria.
6. Final selection of studies for synthesis.

A PRISMA flow diagram should be inserted here once the final database search has been completed and the actual numbers of records identified, screened, excluded, and included have been confirmed.

Figure A1. PRISMA study-selection flow diagram

Figure A1. PRISMA Study-Selection Flow Diagram



Note: Replace all 'TO BE COMPLETED' values with the actual results of the final literature search. The diagram should not contain invented study-selection numbers.

APPENDIX III — DATA EXTRACTION FORM

DATA ITEM	INFORMATION RECORDED
-----------	----------------------

AUTHOR(S)	Name of first author and other authors
YEAR OF PUBLICATION	Publication year
COUNTRY/SETTING	Country and clinical environment
MEDICAL SPECIALTY	Relevant clinical field
STUDY DESIGN	Study methodology
SAMPLE SIZE	Number of patients, cases, images or records
AI TECHNOLOGY	Type of system or model
DIAGNOSTIC TASK	Condition or diagnostic problem assessed
DATA SOURCE	Imaging, pathology, electronic records or other data
COMPARATOR	Clinician, conventional method or standard care
DIAGNOSTIC MEASURES	Sensitivity, specificity, accuracy, AUC and related measures
CLINICAL OUTCOMES	Patient or clinical outcomes were reported
CLINICAL UTILITY	Workflow, decision-making or resource implications
ETHICAL ISSUES	Privacy, fairness, transparency, accountability and autonomy
MAIN FINDINGS	Principal results
STUDY LIMITATIONS	Important methodological limitations
RISK OF BIAS	Overall quality assessment

APPENDIX IV — QUALITY-ASSESSMENT FRAMEWORK

The quality of included studies was assessed according to the characteristics of the evidence being reviewed.

The assessment considered:

- clarity of the research question;

- appropriateness of study design;
- representativeness of the study population;
- quality and completeness of the dataset;
- appropriate separation of training and validation data;
- external validation;
- suitability of the reference standard;
- completeness of outcome reporting;
- statistical methodology;
- transparency of reporting;
- potential sources of bias;
- applicability to routine clinical practice.


For diagnostic accuracy studies, particular attention was given to methodological factors relevant to diagnostic evaluation and reporting. The STARD-AI framework provides specific guidance for reporting AI-centered diagnostic accuracy research, while TRIPOD+AI provides updated guidance for clinical prediction models using regression or machine-learning methods.

APPENDIX V — ADDITIONAL EVIDENCE TABLE

Study	Clinical area	AI application	Main outcome	Principal limitation
Takita et al. (2025)	Clinical diagnosis	Generative AI	Diagnostic performance compared with physicians	Performance varied by model and physician expertise

Shan et al. (2025)	Clinical diagnosis	Large language models	Diagnostic accuracy	High risk of bias in many studies
Digital pathology review (2024)	Pathology	Image analysis	High pooled sensitivity and specificity	Substantial risk of bias/applicability concerns
Wilhelm, Steckelberg and Rebitschek (2025)	Clinical decision-making	Algorithmic systems	Patient-relevant benefits and harms	Limited evidence for patient outcomes
El Arab and Al Moosa (2025)	Multiple specialties	Clinical AI systems	Economic and clinical value	Economic models and reporting varied
Yuan et al. (2025)	Clinical deterioration	Early-warning systems	Patient outcomes	Limited number of prospectively validated studies

APPENDIX VI — CONCEPTUAL FRAMEWORK





Responsible Clinical Use

The framework assumes that diagnostic accuracy is necessary but not sufficient for successful clinical implementation. Clinical value depends on the interaction between technical performance, clinical workflow, professional judgement, patient safety, ethical safeguards, and appropriate governance.